

Greedy Algorithms*

Xiaofeng Gao

Department of Computer Science and Engineering
Shanghai Jiao Tong University, P.R.China

Algorithm Course: Shanghai Jiao Tong University

* Special thanks is given to Prof. Kevin Wayne@Princeton for sharing his lecture notes, and also given to Mr. Hongjian Cao from CS2015@SJTU for producing this slide.

Interval Scheduling: An Introductory Example


- Job j starts at s_j and finishes at f_j .
- Two jobs are **compatible** if they don't overlap.

Goal: find maximum subset of mutually compatible jobs.

Interval Scheduling: An Introductory Example

- Job j starts at s_j and finishes at f_j .
- Two jobs are **compatible** if they don't overlap.

Goal: find maximum subset of mutually compatible jobs.



Optimization Problem: Given a problem Π with domain X , choose a subset or determine a sequence according to some **maximization** or **minimization** objective.

Optimization Problem: Given a problem Π with domain X , choose a subset or determine a sequence according to some **maximization** or **minimization** objective.

General Template: Consider each item x_i in turn from the domain X of problem Π (in some order), make choice that looks **best** at the moment.

Note: it makes a *locally optimal* choice in the hope that this choice will lead to a *globally optimal* solution.

Optimization Problem: Given a problem Π with domain X , choose a subset or determine a sequence according to some **maximization** or **minimization** objective.

General Template: Consider each item x_i in turn from the domain X of problem Π (in some order), make choice that looks **best** at the moment.

*Note: it makes a **locally optimal** choice in the hope that this choice will lead to a **globally optimal** solution.*

Interval Scheduling Problem: Consider jobs in some natural order. Take each job provided to judge its compatibility with the ones already taken.

Attempts

[Earliest start time] Consider jobs in ascending order of s_j .

[Earliest finish time] Consider jobs in ascending order of f_j .

[Shortest interval] Consider jobs in ascending order of $f_j - s_j$.

[Fewest conflicts] For each job j , count the number of conflicting jobs c_j . Schedule in ascending order of c_j .

Attempts

[Earliest start time] Consider jobs in ascending order of s_j .

Counter Example:



[Earliest finish time] Consider jobs in ascending order of f_j .

[Shortest interval] Consider jobs in ascending order of $f_j - s_j$.

[Fewest conflicts] For each job j , count the number of conflicting jobs c_j . Schedule in ascending order of c_j .

[Earliest start time] Consider jobs in ascending order of s_j .

Counter Example:



[Earliest finish time] Consider jobs in ascending order of f_j .

[Shortest interval] Consider jobs in ascending order of $f_j - s_j$.

Counter Example:



[Fewest conflicts] For each job j , count the number of conflicting jobs c_j . Schedule in ascending order of c_j .

[Earliest start time] Consider jobs in ascending order of s_j .



[Earliest finish time] Consider jobs in ascending order of f_j .

[Shortest interval] Consider jobs in ascending order of $f_j - s_j$.



[Fewest conflicts] For each job j , count the number of conflicting jobs c_j . Schedule in ascending order of c_j .



Greedy Interval Scheduling Algorithm

Algorithm 1: Greedy Interval Scheduling

```
1 Sort jobs by finish times so that  $f_1 \leq f_2 \leq \dots \leq f_n$ ;  
2  $A \leftarrow \emptyset$ ; // set of jobs selected  
3 for  $j = 1$  to  $n$  do  
4   if job  $j$  is compatible with  $A$  then  
5      $A \leftarrow A \cup \{j\}$ ;  
6 return  $A$ ;
```

Greedy Interval Scheduling Algorithm

Algorithm 1: Greedy Interval Scheduling

```
1 Sort jobs by finish times so that  $f_1 \leq f_2 \leq \dots \leq f_n$ ;  
2  $A \leftarrow \emptyset$ ; // set of jobs selected  
3 for  $j = 1$  to  $n$  do  
4   if job  $j$  is compatible with  $A$  then  
5      $A \leftarrow A \cup \{j\}$ ;  
6 return  $A$ ;
```

Implementation: $O(n \log n)$.

- After each iteration, set job j^* that was added last to A .
- Job j is compatible with A if $s_j \geq f_{j^*}$.

Theorem. Greedy Interval Scheduling algorithm is optimal.

Theorem. Greedy Interval Scheduling algorithm is optimal.

Proof. (by contradiction) Assume greedy is not optimal.

Theorem. Greedy Interval Scheduling algorithm is optimal.

Proof. (by contradiction) Assume greedy is not optimal.

Let i_1, i_2, \dots, i_k denote set of jobs selected by greedy.

Let j_1, j_2, \dots, j_m denote set of jobs in an optimal solution with $i_1 = j_1$, $i_2 = j_2, \dots, i_r = j_r$ for the largest possible value of r .

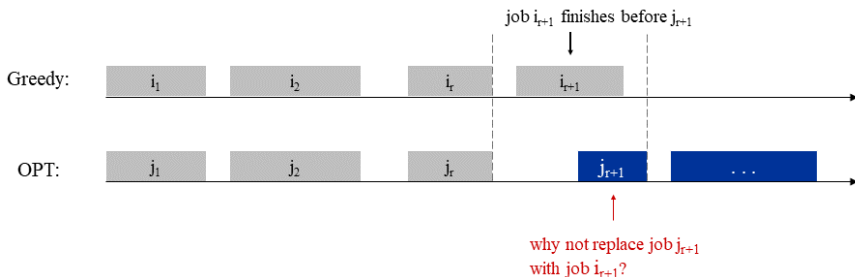
Correctness Analysis

Theorem. Greedy Interval Scheduling algorithm is optimal.

Proof. (by contradiction) Assume greedy is not optimal.

Let i_1, i_2, \dots, i_k denote set of jobs selected by greedy.

Let j_1, j_2, \dots, j_m denote set of jobs in an optimal solution with $i_1 = j_1$, $i_2 = j_2, \dots, i_r = j_r$ for the largest possible value of r .



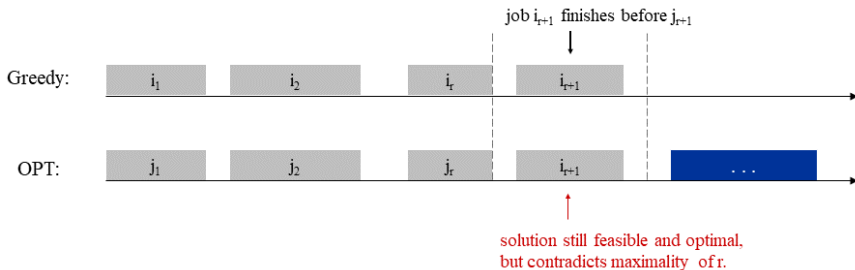
Correctness Analysis

Theorem. Greedy Interval Scheduling algorithm is optimal.

Proof. (by contradiction) Assume greedy is not optimal.

Let i_1, i_2, \dots, i_k denote set of jobs selected by greedy.

Let j_1, j_2, \dots, j_m denote set of jobs in an optimal solution with $i_1 = j_1$, $i_2 = j_2, \dots, i_r = j_r$ for the largest possible value of r .



Interval Partitioning

Lecture j starts at s_j and finishes at f_j .


Goal: find minimum number of classrooms to schedule all lectures so that no two occur at the same time in the same room.

Interval Partitioning

Lecture j starts at s_j and finishes at f_j .

Goal: find minimum number of classrooms to schedule all lectures so that no two occur at the same time in the same room.

Example: This schedule uses 4 classrooms to schedule 10 lectures.




Interval Partitioning

Lecture j starts at s_j and finishes at f_j .

Goal: find minimum number of classrooms to schedule all lectures so that no two occur at the same time in the same room.

Example: This schedule uses only 3.



Lower Bound on Optimal Solution

Definition: The **depth** of a set of open intervals is the maximum number that contain any given time.

Lower Bound on Optimal Solution

Definition: The **depth** of a set of open intervals is the maximum number that contain any given time.


Key observation. Number of classrooms needed \geq depth.

Lower Bound on Optimal Solution

Definition: The **depth** of a set of open intervals is the maximum number that contain any given time.

Key observation. Number of classrooms needed \geq depth.

Example: Depth of schedule = 3 \Rightarrow The schedule is optimal.




Lower Bound on Optimal Solution

Definition: The **depth** of a set of open intervals is the maximum number that contain any given time.

Key observation. Number of classrooms needed \geq depth.

Example: Depth of schedule = 3 \Rightarrow The schedule is optimal.



Question. Does there always exist a schedule equal to depth of intervals?

Greedy Interval Partitioning Algorithm

Algorithm 2: Interval Partitioning Greedy Algorithm

```
1 Sort intervals by starting time so that  $s_1 \leq s_2 \leq \dots \leq s_n$ ;  
2  $d \leftarrow 0$ ;           // number of allocated classrooms  
3 for  $j = 1$  to  $n$  do  
4   if lecture  $j$  is compatible with some classroom  $k$  then  
5     schedule lecture  $j$  in classroom  $k$ ;  
6   else  
7     allocate a new classroom  $d + 1$ ;  
8     schedule lecture  $j$  in classroom  $d + 1$ ;  
9      $d \leftarrow d + 1$ ;  
10 return  $A$ ;
```

Greedy Interval Partitioning Algorithm

Algorithm 2: Interval Partitioning Greedy Algorithm

```
1 Sort intervals by starting time so that  $s_1 \leq s_2 \leq \dots \leq s_n$ ;  
2  $d \leftarrow 0$ ;           // number of allocated classrooms  
3 for  $j = 1$  to  $n$  do  
4   if lecture  $j$  is compatible with some classroom  $k$  then  
5     schedule lecture  $j$  in classroom  $k$ ;  
6   else  
7     allocate a new classroom  $d + 1$ ;  
8     schedule lecture  $j$  in classroom  $d + 1$ ;  
9      $d \leftarrow d + 1$ ;  
10 return  $A$ ;
```

Implementation: $O(n \log n)$.

- For classroom k , maintain the finish time of the last job added.
- Keep the classrooms in a priority queue.

Key Observation. Greedy algorithm never schedules two incompatible lectures in the same classroom.

Key Observation. Greedy algorithm never schedules two incompatible lectures in the same classroom.

Theorem. Greedy algorithm is optimal.

Key Observation. Greedy algorithm never schedules two incompatible lectures in the same classroom.

Theorem. Greedy algorithm is optimal.

Proof. Let d = number of classrooms that the algorithm allocates.

Key Observation. Greedy algorithm never schedules two incompatible lectures in the same classroom.

Theorem. Greedy algorithm is optimal.

Proof. Let d = number of classrooms that the algorithm allocates.

Classroom d is opened because we needed to schedule a job, say j , that is incompatible with all $d - 1$ other classrooms. (These d jobs each end after s_j .)

Key Observation. Greedy algorithm never schedules two incompatible lectures in the same classroom.

Theorem. Greedy algorithm is optimal.

Proof. Let d = number of classrooms that the algorithm allocates.

Classroom d is opened because we needed to schedule a job, say j , that is incompatible with all $d - 1$ other classrooms. (These d jobs each end after s_j .)

Since we sorted by start time, all these incompatibilities are caused by lectures that start no later than s_j . Thus, we have d lectures overlapping at time $s_j + \varepsilon$. □

Key Observation. Greedy algorithm never schedules two incompatible lectures in the same classroom.

Theorem. Greedy algorithm is optimal.

Proof. Let d = number of classrooms that the algorithm allocates.

Classroom d is opened because we needed to schedule a job, say j , that is incompatible with all $d - 1$ other classrooms. (These d jobs each end after s_j .)

Since we sorted by start time, all these incompatibilities are caused by lectures that start no later than s_j . Thus, we have d lectures overlapping at time $s_j + \varepsilon$. □

Key observation \Rightarrow all schedules use $\geq d$ classrooms.

Scheduling to Minimize Lateness


- Single resource processes one job at a time.
- Job j requires t_j units of processing time and is due at time d_j .
- If j starts at time s_j , it finishes at time $f_j = s_j + t_j$.
- Lateness: $l_j = \max\{0, f_j - d_j\}$.

Goal: schedule all jobs to minimize maximum lateness $L = \max l_j$.

Scheduling to Minimize Lateness

- Single resource processes one job at a time.
- Job j requires t_j units of processing time and is due at time d_j .
- If j starts at time s_j , it finishes at time $f_j = s_j + t_j$.
- Lateness: $l_j = \max\{0, f_j - d_j\}$.

Goal: schedule all jobs to minimize maximum lateness $L = \max l_j$.



Attempt: Consider jobs in ascending order by some strategy

[Shortest processing time first] Sort by processing time t_j .

[Earliest deadline first] Sort by deadline d_j .

[Smallest slack] Sort by slack $d_j - t_j$.

Attempt: Consider jobs in ascending order by some strategy

[Shortest processing time first] Sort by processing time t_j .

	1	2
t_j	1	10
d_j	100	10

counterexample

[Earliest deadline first] Sort by deadline d_j .

[Smallest slack] Sort by slack $d_j - t_j$.

Attempt: Consider jobs in ascending order by some strategy

[Shortest processing time first] Sort by processing time t_j .

	1	2
t_j	1	10
d_j	100	10

counterexample

[Earliest deadline first] Sort by deadline d_j .

[Smallest slack] Sort by slack $d_j - t_j$.

	1	2
t_j	1	10
d_j	2	10

counterexample

A Greedy Algorithm: Earliest Deadline First

Algorithm 3: Greedy Minimizing Lateness


- 1 Sort n jobs by deadline so that $d_1 \leq d_2 \leq \dots \leq d_n$;
 - 2 $t \leftarrow 0$;
 - 3 **for** $j = 1$ **to** n **do**
 - 4 Assign job j to interval $[t, t + t_j]$;
 - 5 $s_j \leftarrow t, f_j \leftarrow t + t_j$;
 - 6 $t \leftarrow t + t_j$;
 - 7 **return** intervals $[s_j, f_j]$;
-

A Greedy Algorithm: Earliest Deadline First

Algorithm 3: Greedy Minimizing Lateness


- 1 Sort n jobs by deadline so that $d_1 \leq d_2 \leq \dots \leq d_n$;
 - 2 $t \leftarrow 0$;
 - 3 **for** $j = 1$ **to** n **do**
 - 4 Assign job j to interval $[t, t + t_j]$;
 - 5 $s_j \leftarrow t, f_j \leftarrow t + t_j$;
 - 6 $t \leftarrow t + t_j$;
 - 7 **return** intervals $[s_j, f_j]$;
-

Implementation: $O(n \log n)$.




Correctness Proof: Reduce Optimal Solution

Observation. There exists an optimal schedule with no **idle time**.



Correctness Proof: Reduce Optimal Solution

Observation. There exists an optimal schedule with no **idle time**.



Observation. The greedy schedule has no idle time.

Correctness Proof: Optimal Solution vs Algorithm Solution

Definition. Given a schedule S , an inversion is a pair of jobs i and j such that: $i < j$ but j scheduled before i .



[as before, we assume jobs are numbered so that $d_1 \leq d_2 \leq \dots \leq d_n$]

Correctness Proof: Optimal Solution vs Algorithm Solution

Definition. Given a schedule S , an inversion is a pair of jobs i and j such that: $i < j$ but j scheduled before i .




[as before, we assume jobs are numbered so that $d_1 \leq d_2 \leq \dots \leq d_n$]

Observation. Greedy schedule has no inversions.

Observation. If a schedule (with no idle time) has an inversion, it has one with a pair of inverted jobs scheduled consecutively.


Correctness Proof: Inversions

Claim. Swapping two consecutive, inverted jobs reduces the number of inversions by one and does not increase the max lateness.



Correctness Proof: Inversions

Claim. Swapping two consecutive, inverted jobs reduces the number of inversions by one and does not increase the max lateness.




Proof. Let l be the lateness before the swap, and let l' be it afterwards.

- $l'_k = l_k$ for all $k \neq i, j$
- $l'_i \leq l_i$
- If job j is late:

Correctness Proof: Inversions

Claim. Swapping two consecutive, inverted jobs reduces the number of inversions by one and does not increase the max lateness.



Proof. Let l be the lateness before the swap, and let l' be it afterwards.

- $l'_k = l_k$ for all $k \neq i, j$
- $l'_j = f'_j - d_j$ (definition)
 $= f_i - d_j$ (j finishes at time f_i)
- $l'_i \leq l_i$
- If job j is late:
 $\leq f_i - d_i$ ($i < j$)
 $\leq l_i$ (definition)

Theorem. Greedy schedule S is optimal.

Theorem. Greedy schedule S is optimal.

Proof. Define S^* to be an optimal schedule that has the fewest number of inversions. (We can assume S^* has no idle time.)

Theorem. Greedy schedule S is optimal.

Proof. Define S^* to be an optimal schedule that has the fewest number of inversions. (We can assume S^* has no idle time.)

- If S^* has no inversions, then $S = S^*$.

Theorem. Greedy schedule S is optimal.

Proof. Define S^* to be an optimal schedule that has the fewest number of inversions. (We can assume S^* has no idle time.)

- If S^* has no inversions, then $S = S^*$.
- If S^* has an inversion, let $i - j$ be an adjacent inversion.

Theorem. Greedy schedule S is optimal.

Proof. Define S^* to be an optimal schedule that has the fewest number of inversions. (We can assume S^* has no idle time.)

- If S^* has no inversions, then $S = S^*$.
- If S^* has an inversion, let $i - j$ be an adjacent inversion.

Swapping i and j does not increase the maximum lateness and strictly decreases the number of inversions.

Theorem. Greedy schedule S is optimal.

Proof. Define S^* to be an optimal schedule that has the fewest number of inversions. (We can assume S^* has no idle time.)

- If S^* has no inversions, then $S = S^*$.
- If S^* has an inversion, let $i - j$ be an adjacent inversion.

Swapping i and j does not increase the maximum lateness and strictly decreases the number of inversions. This contradicts definition of S^* . □

Greedy algorithm stays ahead. Show that after each step of the greedy algorithm, its solution is at least as good as any other algorithm's.

Structural. Discover a simple "structural" bound asserting that every possible solution must have a certain value. Then show that your algorithm always achieves this bound.

Exchange argument. Gradually transform any solution to the one found by the greedy algorithm without hurting its quality.

Other greedy algorithms. Kruskal, Prim, Dijkstra, Huffman, ...

Coin Changing

Goal. Given US currency denominations:

1 (cent), 5 (nickel), 10 (dime), 25 (quarter), 100 (dollar),

devise a changing method using fewest number of coins.

Coin Changing

Goal. Given US currency denominations:

1 (cent), 5 (nickel), 10 (dime), 25 (quarter), 100 (dollar),

devise a changing method using fewest number of coins.

Example. 34¢.



Coin Changing

Goal. Given US currency denominations:

1 (cent), 5 (nickel), 10 (dime), 25 (quarter), 100 (dollar),

devise a changing method using fewest number of coins.

Example. 34¢.



Cashier's algorithm. At each iteration, add coin of the largest value that does not take us past the amount to be paid.

Coin Changing

Goal. Given US currency denominations:

1 (cent), 5 (nickel), 10 (dime), 25 (quarter), 100 (dollar),

devise a changing method using fewest number of coins.

Example. 34¢.



Cashier's algorithm. At each iteration, add coin of the largest value that does not take us past the amount to be paid.

Example. \$2.89.



Algorithm 4: Cashier's Algorithm

```
1 Sort coins denominations by value:  $c_1 < c_2 < \dots < c_n$ ;  
2  $S \leftarrow \emptyset$ ; // coins selected  
3 while  $x \neq 0$  do  
4   let  $k$  be largest integer such that  $c_k < x$ ;  
5   if  $k = 0$  then  
6     return "no solution found";  
7    $x \leftarrow x - c_k$ ;  
8    $S \leftarrow S \cup \{k\}$ ;  
9 return  $S$ ;
```

Algorithm 4: Cashier's Algorithm

```
1 Sort coins denominations by value:  $c_1 < c_2 < \dots < c_n$ ;  
2  $S \leftarrow \emptyset$ ; // coins selected  
3 while  $x \neq 0$  do  
4   let  $k$  be largest integer such that  $c_k < x$ ;  
5   if  $k = 0$  then  
6     return "no solution found";  
7    $x \leftarrow x - c_k$ ;  
8    $S \leftarrow S \cup \{k\}$ ;  
9 return  $S$ ;
```

Question. Is cashier's algorithm optimal?

Properties of Optimal Solution

Property. Number of pennies ≤ 4 .

Proof. Replace 5 pennies with 1 nickel.

Properties of Optimal Solution

Property. Number of pennies ≤ 4 .

Proof. Replace 5 pennies with 1 nickel.

Property. Number of nickels ≤ 1 .

Property. Number of quarters ≤ 3 .

Properties of Optimal Solution

Property. Number of pennies ≤ 4 .

Proof. Replace 5 pennies with 1 nickel.

Property. Number of nickels ≤ 1 .

Property. Number of quarters ≤ 3 .

Property. Number of pennies ≤ 2 .

Proof.

- Replace 3 dimes and 0 nickels with 1 quarter and 1 nickel;
- Replace 2 dimes and 1 nickel with 1 quarter.
- Recall: at most 1 nickel.



Theorem. Greedy algorithm is optimal for U.S. coinage.

Theorem. Greedy algorithm is optimal for U.S. coinage.

Proof. (by induction on x) Consider an optimal way to change $c_k < x < c_{k+1}$: greedy takes coin k . We claim that any optimal solution must also take coin k .

Theorem. Greedy algorithm is optimal for U.S. coinage.

Proof. (by induction on x) Consider an optimal way to change $c_k < x < c_{k+1}$: greedy takes coin k . We claim that any optimal solution must also take coin k .

If not, it needs enough coins of type c_1, \dots, c_{k-1} to add up to x .

k	c_k	All optimal solutions must satisfy	Max value of coins 1, 2, ..., $k-1$ in any OPT
1	1	$P \leq 4$	-
2	5	$N \leq 1$	4
3	10	$N + D \leq 2$	$4 + 5 = 9$
4	25	$Q \leq 3$	$20 + 4 = 24$
5	100	no limit	$75 + 24 = 99$

Theorem. Greedy algorithm is optimal for U.S. coinage.

Proof. (by induction on x) Consider an optimal way to change $c_k < x < c_{k+1}$: greedy takes coin k . We claim that any optimal solution must also take coin k .

If not, it needs enough coins of type c_1, \dots, c_{k-1} to add up to x .

k	c_k	All optimal solutions must satisfy	Max value of coins 1, 2, ..., $k-1$ in any OPT
1	1	$P \leq 4$	-
2	5	$N \leq 1$	4
3	10	$N + D \leq 2$	$4 + 5 = 9$
4	25	$Q \leq 3$	$20 + 4 = 24$
5	100	no limit	$75 + 24 = 99$

Problem reduces to coin-changing $x - c_k$ cents, which, by induction, is optimally solved by greedy algorithm.

Is Cashier's Algorithm Work for Any Denominations?

Observation 1. Greedy is sub-optimal for US postal denominations:

1, 10, 21, 34, 70, 100, 350, 1225, 1500.

Is Cashier's Algorithm Work for Any Denominations?

Observation 1. Greedy is sub-optimal for US postal denominations:

1, 10, 21, 34, 70, 100, 350, 1225, 1500.

Counter example. 140¢. (Greedy: 100, 34, six 1's; Optimal: 70, 70.)



Is Cashier's Algorithm Work for Any Denominations?

Observation 1. Greedy is sub-optimal for US postal denominations:

1, 10, 21, 34, 70, 100, 350, 1225, 1500.


Counter example. 140¢. (Greedy: 100, 34, six 1's; Optimal: 70, 70.)



Observation 2. Even no feasible solution with system $\mathcal{C}=\{7, 8, 9\}$.

- Cashier's algorithm: $15\text{¢} = 9 + \text{???}$
- Optimal: $15\text{¢} = 7 + 8$.

Movie: Wall Street (1987)



†

Greed is good.

Greed is right.

Greed works.

Greed clarifies, cuts through, and captures the essence of the evolutionary spirit.

— Gordon Gecko
(Michael Douglas)

† Watch the movie segment at the class webpage.